

Homework 1: Binary Choice

Due date: 8.11

The goal of this homework is to introduce you to google colab notebooks and simple binary choice problems.

1. Go over the google colab [introduction notebook](#) and see that you understand the basic components of google colaboratory.
2. Go over the introduction notebook provided and see that you understand the basic components of the training pipeline: user model, population model, training models and metrics. Make sure you get an accuracy higher than 80%.
3. Create the **SigmoidUserModel(UserModel)** class. Like the NoiseUserModel from the introduction notebook, this class models user valuations on items.

The valuation of a user u on item x is modeled as follows:

$$v(u, x) = \text{sigmoid}(u^T x) - \frac{1}{2}$$
$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

Also, create the SigmoidPopulationModel(PopulationModel) accordingly.

Note: use `np.random.randn(num_features)` to generate user vectors u , similarly to NoiseUserModel.

4. Create a dataset with the `generate_synth_data(num_features, num_items)` function and the SigmoidPopulationModel class you created in section (3). Then run it through a training pipeline with the accuracy metric and the LinearRegression() training model. Show the achieved accuracy.
5. Repeat section (4) but with a logistic regression training model.
Note: replace `LinearRegression()` with `MultiOutputClassifier(LogisticRegression())`.
6. Run 100 iterations of sections (4)+(5). Meaning in each iterations generate a new dataset with SigmoidPopulationModel and train once with a linear regression model and once with a logistic regression model. Plot a graph of the accuracies in each iteration (2 graphs in the same figure – one for each training model accuracies).
7. Explain the results you got in section (6). Why do you think the graphs look like that?
8. Create the `welfare(training_model, population_model, model_args, X)` evaluation metric which will be used in our next workshop.

The welfare of users U with respect to items M is defined as follows:

$$Welfare = \frac{1}{|U|} \sum_{u \in U} \sum_{m \in M} \mathbb{I}\{\text{sign}(f(u, m)) = 1\} \cdot v(u, m)$$

Where $f(u, m)$ is the predicted valuation of the prediction model of user u on item m and $v(u, m)$ is the true valuation of user u on item m . \mathbb{I} is the indicator function

$$(I(x) = \begin{cases} 1, & x = \text{True} \\ 0, & x = \text{False} \end{cases}).$$

9. Create the `MSE(training_model, population_model, model_args, X)` evaluation metric which will be used in our next workshop.

The MSE (mean squared error) of users U with respect to items M is defined as follows:

$$MSE = \frac{1}{|U||M|} \sum_{u \in U} \sum_{m \in M} (f(u, m) - v(u, m))^2$$

Submission guidelines:

1. You can (if you want) create a copy of the introduction notebook and fill in your solutions at the end of it.
2. Submit a .ipynb file with your solution.
3. Add a text block in the beginning of your notebook with your IDs.
4. Indicate clearly with a text block the sections of your solutions
5. Answer dry questions in text blocks and wet questions in code blocks.
6. For any questions regarding this homework, contact [Sagi](#).

Example for submission format:

IDs:

- 123456789
- 987654321

Question 1

```
[ ] # My wet solution to question 1
```

Question 2

```
[ ] # My wet solution to question 2
```

Maybe some dry explanations

Question 3

```
[ ] # My wet solution to question 3
```

```
[ ] # Continuing my wet solution to question 3
```